

Retail Forecasting Project

ETC3550

Chelaka Paranaheewa


Statistical features of Australian Retail

Table 1: A few rows of the dataset, Other Retailing in South Australia

State	Industry	Series ID	Month	Turnover
South Australia	Other retailing	A3349433W	1982 Apr	34.2
South Australia	Other retailing	A3349433W	1982 May	34.4
South Australia	Other retailing	A3349433W	1982 Jun	32.7
South Australia	Other retailing	A3349433W	1982 Jul	36.2
South Australia	Other retailing	A3349433W	1982 Aug	36.1
South Australia	Other retailing	A3349433W	1982 Sep	36.0


Australia Retail Turnover

Original



Plotting the original dataset, we can see a general trend upwards. The dataset also has a seasonal pattern which during the early 1980 - 1990, was relatively small compared to later years where the spike grows to large proportions.

Seasonal Plot: Australia Retail Turnover



From the seasonal plot, there is a clear increase in retail sales during the month December which suggests that there is strong seasonality with the data.

Transformations and Differencings

Transformations

To prepare the dataset for ARIMA modelling, the data needs to be transformed so that the variance across the dataset remains relatively constant. As mentioned before the original dataset does not have a constant seasonal variation, since the beginning has small peaks in the seasonal variation which gradually grow over the course of approximately two decades.

To stabilise the seasonal variation, a Box-Cox transformation can be used. A Box-Cox transformation needs a value for λ to transform the data. A $\lambda = 0$ will perform a logarithmic transformation to the data, whereas $\lambda \neq 0$ will perform an exponential transformation.

```
lambda <- myseries %>%
  features(Turnover, features = guerrero) %>%
  pull(lambda_guerrero)

myseries %>% autoplot(box_cox(Turnover, lambda)) +
  labs(
    title = "Australia Retail Turnover",
```

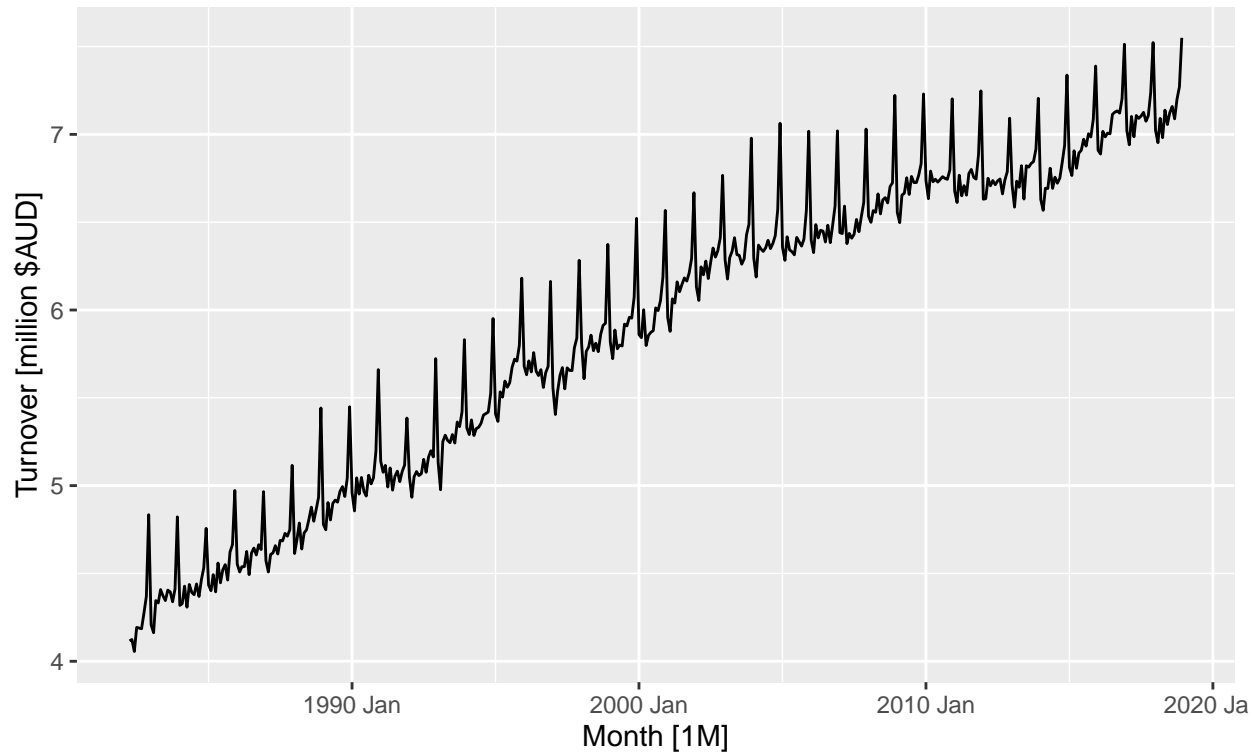
```

    subtitle = "Box-Cox transformed (Guerrero's method)"
  ) +
  ylab("Turnover [million $AUD]")

```

Australia Retail Turnover

Box-Cox transformed (Guerrero's method)



Using a λ value close to 0 such as 0.0845031 gives the best transformation that keeps the seasonal variations constant through out the time series. This value was determined by using Guerrero's method and then manually checking values around it to verify its the best value for λ .

Differencing

To objectively determine if the dataset needs differencing we will use `unitroot_kpss` & `unitroot_nsdiffs`.

```

myseries %>%
  features(Turnover, unitroot_kpss) %>%
  kable(caption = "Kwiatkowski-Phillips-Schmidt-Shin Test")

```

Table 2: Kwiatkowski-Phillips-Schmidt-Shin Test

State	Industry	kpss_stat	kpss_pvalue
South Australia	Other retailing	7.376257	0.01

```

myseries %>%
  features(difference(Turnover, 12), unitroot_kpss) %>%
  kable(caption = "Kwiatkowski-Phillips-Schmidt-Shin Test")

```

Table 3: Kwiatkowski-Phillips-Schmidt-Shin Test

State	Industry	kpss_stat	kpss_pvalue
South Australia	Other retailing	0.3540707	0.0969523

The KPSS test is performing a hypothesis test to verify that the data is stationary. However, the null hypothesis is rejected since the `kpss_pvalue` is less than 0.05, hence indicating that the data is not stationary and needs differencing to make it stationary.

After a first order differencing, the hypothesis test gives a p value of 0.10 which is greater than 0.05 which means the null hypothesis is not rejected. Thus making the first difference of the dataset, stationary.

Based on Kwiatkowski-Phillips-Schmidt-Shin's Test, we have determined that the data needs differencing to make it stationary, while we tested with a first stage differencing we need to verify this using the `unitroot_ndiffs` which reveals the number of differencing that is needed.

```
myseries %>%
  features(box_cox(Turnover, lambda), unitroot_ndiffs) %>%
  kable(caption = "Number of differences required for a stationary series")
```

Table 4: Number of differences required for a stationary series

State	Industry	ndiffs
South Australia	Other retailing	1

According to `unitroot_ndiffs` the data only needs to perform a first stage differencing.

```
myseries %>%
  features(box_cox(Turnover, lambda), unitroot_nsdiffs) %>%
  kable(
    caption =
      "Number of seasonal differences required for a stationary series"
  )
```

Table 5: Number of seasonal differences required for a stationary series

State	Industry	nsdiffs
South Australia	Other retailing	1

According to `unitroot_ndiffs` the data only needs to perform a first stage seasonal differencing.

Modelling ARIMA and ETS models

ETS Modelling

For the ETS models, the best method to short list possible candidate models is to use the AIC values that are outputted after models are trained. AIC, otherwise known as Akaike's Information Criterion, is defined as $AIC = T \log(\frac{SSE}{T}) + 2(k + 2)$. Using it to find a model by minimising the AIC will result in a model that is good at forecasting.

To find the model with the lowest AIC value, we need to check all the different models which can be done by `model(ETS(Turnover))` or manually testing all combinations.

```

progressr::with_progress(
  trained_etsmodel <- training %>%
    model(
      # Additive
      ANN = ETS(Turnover ~ error("A") + trend("N") + season("N")),
      ANA = ETS(Turnover ~ error("A") + trend("N") + season("A")),
      ANM = ETS(Turnover ~ error("A") + trend("N") + season("M")),
      AAN = ETS(Turnover ~ error("A") + trend("A") + season("N")),
      AAA = ETS(Turnover ~ error("A") + trend("A") + season("A")),
      AAM = ETS(Turnover ~ error("A") + trend("A") + season("M")),
      AAdN = ETS(Turnover ~ error("A") + trend("Ad") + season("N")),
      AAdA = ETS(Turnover ~ error("A") + trend("Ad") + season("A")),
      AAdM = ETS(Turnover ~ error("A") + trend("Ad") + season("M")),

      # Multiplicative
      MNN = ETS(Turnover ~ error("M") + trend("N") + season("N")),
      MNA = ETS(Turnover ~ error("M") + trend("N") + season("A")),
      MNM = ETS(Turnover ~ error("M") + trend("N") + season("M")),
      MAN = ETS(Turnover ~ error("M") + trend("A") + season("N")),
      MAA = ETS(Turnover ~ error("M") + trend("A") + season("A")),
      MAM = ETS(Turnover ~ error("M") + trend("A") + season("M")), #*
      MAdN = ETS(Turnover ~ error("M") + trend("Ad") + season("N")),
      MAdA = ETS(Turnover ~ error("M") + trend("Ad") + season("A")),
      MAdM = ETS(Turnover ~ error("M") + trend("Ad") + season("M"))
    )
)

trained_etsmodel %>%
  glance() %>%
  select(.model, AIC, AICc) %>%
  arrange(AIC) %>%
  kable(caption = "ETS Models ranked based on AIC")

```

Table 6: ETS Models ranked based on AIC

.model	AIC	AICc
MAM	3708.183	3709.717
MAdM	3720.894	3722.612
MNM	3837.702	3838.899
AAM	3904.291	3905.825
AAdM	3906.502	3908.221
ANM	3943.362	3944.559
ANA	4034.415	4035.612
MAA	4038.320	4039.854
MNA	4046.030	4047.227
AAA	4058.046	4059.579
MAdA	4066.866	4068.585
AAdA	4079.408	4081.127
MAN	4737.305	4737.451
MAdN	4748.376	4748.580
MNN	4758.211	4758.269
AAdN	4930.541	4930.746
AAN	4931.342	4931.488


.model	AIC	AICc
ANN	4950.265	4950.323

By testing all combinations we see that the top 6 ETS models are: MAM, MAdM, MNM, AAM, AAdM, ANM. The list of ETS models can be further reduced by plotting a time series decomposition which shows the relationship between the seasonality and error components.

```
training %>%
  model(STL(
    # box_cox(Turnover, lambda) ~
    Turnover ~
      trend(window = 13) + season(window = "periodic"),
    robust = TRUE
  )) %>%
  components() %>%
  autoplot()
```

STL decomposition

Turnover = trend + season_year + remainder



The decomposition shown in the graph indicates that the time series is multiplicative. Since the transformation done to the dataset was a box_cox transformation using a lambda value of 0.0845031 which is closer to 0. And according to the box_cox transformation lambda value being zero is a log transformation. Hence the multiplicative decomposition.

So it can be deduced that the time series has a multiplicative relationship both the seasonality and error components which means the final short list of ETS models are: MNM, MAM, MAdM.

After short listing, the best ETS model according to the AIC values is a MAM with a AIC value of 3709.7172205

```

bind_rows(
  trained_etsmodel %>%
    accuracy(),
  trained_etsmodel %>%
    forecast(h = "2 years") %>%
    accuracy(myseries)
) %>%
  arrange(MASE) %>%
  select(-State, -Industry, -ME, -MPE, -ACF1) %>%
  kable(caption = "Shortlist of the best ETS Models ranked on MASE")

```

Table 7: Shortlist of the best ETS Models ranked on MASE

.model	.type	RMSE	MAE	MAPE	MASE	RMSSE
MAM	Training	5.119750	3.631144	2.852607	0.4308353	0.4713354
MAdM	Training	5.210748	3.683136	2.884265	0.4370042	0.4797128
MNM	Training	5.559450	4.072462	3.294554	0.4831977	0.5118151
MNM	Test	6.722146	4.839242	1.812264	0.5741762	0.6188555
MAM	Test	8.277969	6.249165	2.226650	0.7414636	0.7620879
MAdM	Test	8.230045	6.369881	2.279294	0.7557866	0.7576759

After testing the models with the last 2 years of the dataset, the MNM model proves to be the best across RMSE, MAPE and MASE. The MAM and MAdM are also just as good with both tied for second place since their RMSE, MAPE and MASE values are marginally different.


ARIMA Modelling

From section Transformations and Differencings, we know that the dataset needs to be transformed and differenced once to make the data stationary. So we can directly plot the ACF and PACF plots of the stationary dataset.

```

training %>% gg_tsdisplay(
  difference(box_cox(Turnover, lambda), 12),
  plot_type = "partial", lag = 64
) + labs(title = "ACF and PACF plots of transformed and differenced Turnover")

```



The ACF plot shows a decaying sinusoidal pattern and by looking at the PACF the last significant lag is at lag 3 which may suggest a non-seasonal AR(3).

There is exponential decay in the seasonal lags of the PACF which could indicate a seasonal MA(3) component.

From the ACF and PACF plots we have a possible ARIMA model to start from $\text{ARIMA}(\text{box_cox}(\text{Turnover}, \lambda) \sim 1 + \text{pdq}(3, 0, 0) + \text{PDQ}(0, 1, 3))$

To get better values we could make the p, P, q & Q values drift by 1 value and also checking with the addition of a constant.

```
progressr::with_progress(
  trained_arimamodel <- training %>% model(
    ARIMA1300013 = ARIMA(box_cox(Turnover, lambda) ~ 1 + pdq(3, 0, 0) + PDQ(0, 1, 3)),
    ARIMA0300013 = ARIMA(box_cox(Turnover, lambda) ~ 0 + pdq(3, 0, 0) + PDQ(0, 1, 3)),
    ARIMA1200013 = ARIMA(box_cox(Turnover, lambda) ~ 1 + pdq(2, 0, 0) + PDQ(0, 1, 3)),
    ARIMA0200013 = ARIMA(box_cox(Turnover, lambda) ~ 0 + pdq(2, 0, 0) + PDQ(0, 1, 3)),
    ARIMA1201013 = ARIMA(box_cox(Turnover, lambda) ~ 1 + pdq(2, 0, 1) + PDQ(0, 1, 3)),
    ARIMA0201013 = ARIMA(box_cox(Turnover, lambda) ~ 0 + pdq(2, 0, 1) + PDQ(0, 1, 3)),
    ARIMA1300112 = ARIMA(box_cox(Turnover, lambda) ~ 1 + pdq(3, 0, 0) + PDQ(1, 1, 2)),
    ARIMA0300112 = ARIMA(box_cox(Turnover, lambda) ~ 0 + pdq(3, 0, 0) + PDQ(1, 1, 2)),
    ARIMA1200112 = ARIMA(box_cox(Turnover, lambda) ~ 1 + pdq(2, 0, 0) + PDQ(1, 1, 2)),
    ARIMA0200112 = ARIMA(box_cox(Turnover, lambda) ~ 0 + pdq(2, 0, 0) + PDQ(1, 1, 2))
  )
)

trained_arimamodel %>%
  glance() %>%
```



```
select(.model, AIC, AICc) %>%
  arrange(AIC) %>%
  kable(caption = "ARIMA Models ranked based on AIC")
```

Table 8: ARIMA Models ranked based on AIC

.model	AIC	AICc
ARIMA1300013	-1173.091	-1172.728
ARIMA1300112	-1172.961	-1172.597
ARIMA1201013	-1169.124	-1168.760
ARIMA0300013	-1168.133	-1167.851
ARIMA0300112	-1168.082	-1167.800
ARIMA0201013	-1165.399	-1165.116
ARIMA1200013	-1141.190	-1140.908
ARIMA1200112	-1141.095	-1140.813
ARIMA0200013	-1131.407	-1131.196
ARIMA0200112	-1131.134	-1130.923

By checking slightly different variations of the initially obtained ARIMA model we can tell that ARIMA1300013 and ARIMA1300112 fit the data the best since their AIC values are very low and only 0.13 apart.

```
bind_rows(
  trained_arimamodel %>%
    accuracy(),
  trained_arimamodel %>%
    forecast(h = "2 years") %>%
    accuracy(myseries)
) %>%
  arrange(MASE) %>%
  select(-State, -Industry, -ME, -MPE, -ACF1) %>%
  kable(caption = "Shortlist of the best ARIMA Models ranked on MASE")
```

Table 9: Shortlist of the best ARIMA Models ranked on MASE

.model	.type	RMSE	MAE	MAPE	MASE	RMSSE
ARIMA1201013	Training	4.952779	3.563100	2.793952	0.4227619	0.4559636
ARIMA1200112	Training	5.038566	3.580408	2.847238	0.4248155	0.4638614
ARIMA1200013	Training	5.037117	3.580471	2.848343	0.4248230	0.4637279
ARIMA1300013	Training	5.048520	3.582733	2.765884	0.4250913	0.4647777
ARIMA1300112	Training	5.051571	3.584964	2.766082	0.4253561	0.4650586
ARIMA0201013	Training	4.992002	3.600041	2.823828	0.4271449	0.4595746
ARIMA0300112	Training	5.080727	3.640988	2.812820	0.4320033	0.4677428
ARIMA0300013	Training	5.086617	3.646139	2.817608	0.4326145	0.4682851
ARIMA0200013	Training	5.064571	3.652992	2.916282	0.4334276	0.4662554
ARIMA0200112	Training	5.091195	3.678640	2.937793	0.4364707	0.4687065
ARIMA0200112	Test	10.199216	7.534512	2.753414	0.8939701	0.9389621
ARIMA0200013	Test	11.490382	8.632952	3.156135	1.0242999	1.0578296
ARIMA0300112	Test	13.848926	10.772215	3.953826	1.2781236	1.2749624
ARIMA0300013	Test	13.822333	10.772470	3.955113	1.2781539	1.2725141
ARIMA0201013	Test	16.329307	13.299378	4.896372	1.5779715	1.5033116
ARIMA1201013	Test	20.765357	17.286884	6.366140	2.0510893	1.9117041
ARIMA1300112	Test	22.132087	18.125307	6.659770	2.1505682	2.0375282

.model	.type	RMSE	MAE	MAPE	MASE	RMSSE
ARIMA1300013	Test	22.341353	18.350703	6.745479	2.1773114	2.0567937
ARIMA1200112	Test	28.183447	23.443429	8.613578	2.7815634	2.5946296
ARIMA1200013	Test	28.354257	23.663139	8.697970	2.8076321	2.6103548

After testing the models with the last 2 years of the dataset, the ARIMA1300013 and ARIMA1300112 models performed the did not perform the best during the testing phase.

Best Models for Forecasting

Best ETS Model

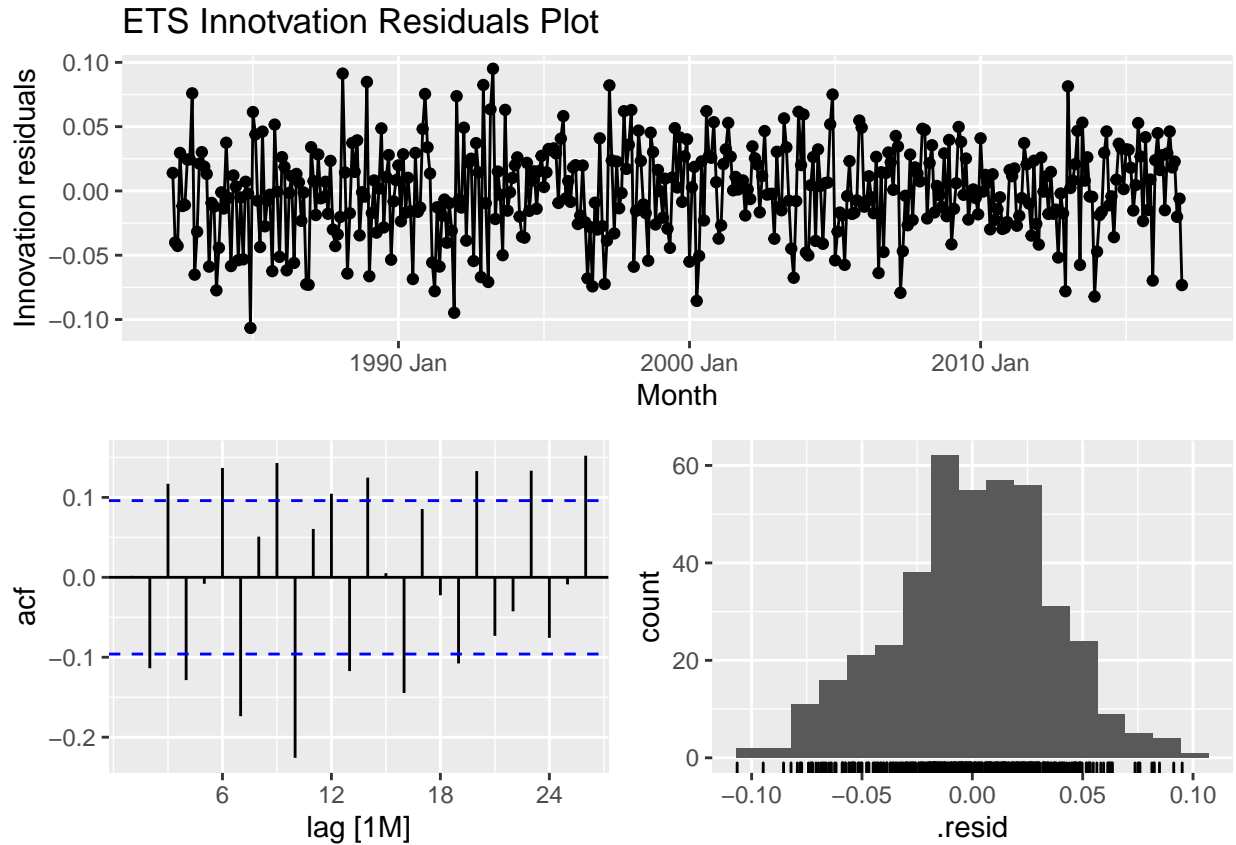
```
MAM <- trained_etsmodel %>%
  select(MAM)

MAM %>%
  report()

## Series: Turnover
## Model: ETS(M,A,M)
## Smoothing parameters:
##   alpha = 0.4954101
##   beta  = 0.0001003904
##   gamma = 0.0001120509
##
## Initial states:
##   l[0]    b[0]    s[0]    s[-1]    s[-2]    s[-3]    s[-4]    s[-5]
## 35.40998 0.4843682 0.9788309 0.8989234 0.9427522 1.353439 1.041756 0.9953029
##   s[-6]    s[-7]    s[-8]    s[-9]    s[-10]    s[-11]
## 0.9737481 0.9843161 0.9740065 0.9389047 0.9784049 0.9396147
##
## sigma^2: 0.0013
##
##      AIC      AICc      BIC
## 3708.183 3709.717 3776.746
```

The best ETS model was the MAM model since it had the lowest AIC value and during the testing set it also had a low MASE score which was not the lowest but relatively close to the best.

```
MAM %>%
  gg_tsresiduals() + labs(title = "ETS Innotvation Residuals Plot")
```



```
MAM %>%
  augment() %>%
  features(.innov, ljung_box, lag = 24, dof = 16) %>%
  kable(caption = "MAM Ljung Box Test")
```

Table 10: MAM Ljung Box Test

.model	lb_stat	lb_pvalue
MAM	128.8372	0

From the innovation residuals plot and its corresponding ACF plot shows that there might still be some correlation that was missed. Which means that model is not perfect and will not provide the most accurate forecasts. With that being said, the the true value will still land with in its forecast interval.

```
MAM_fc <- MAM %>%
  forecast(h = "2 years")

MAM_fc %>%
  hilo(level = c(80, 95)) %>%
  select(-.model, -Turnover) %>%
  kable(caption = "ETS Point Forecast and Point Interval (80% & 95%)")
```

Table 11: ETS Point Forecast and Point Interval (80% & 95%)


Month	.mean	80%	95%
2017 Jan	243.8380	[232.4882, 255.1878]80	[226.4800, 261.1960]95
2017 Feb	232.9391	[220.8415, 245.0368]80	[214.4374, 251.4409]95
2017 Mar	254.1250	[239.6907, 268.5594]80	[232.0496, 276.2005]95
2017 Apr	244.4036	[229.4295, 259.3776]80	[221.5028, 267.3044]95
2017 May	254.9684	[238.2887, 271.6481]80	[229.4590, 280.4777]95
2017 Jun	245.1350	[228.1446, 262.1253]80	[219.1505, 271.1194]95
2017 Jul	254.7704	[236.1763, 273.3645]80	[226.3332, 283.2077]95
2017 Aug	257.9484	[238.2228, 277.6740]80	[227.7807, 288.1161]95
2017 Sep	255.6572	[235.2567, 276.0578]80	[224.4573, 286.8572]95
2017 Oct	261.8057	[240.0811, 283.5302]80	[228.5808, 295.0305]95
2017 Nov	274.5335	[250.9137, 298.1534]80	[238.4101, 310.6570]95
2017 Dec	357.3301	[325.5349, 389.1252]80	[308.7036, 405.9565]95
2018 Jan	249.3663	[226.4690, 272.2636]80	[214.3479, 284.3847]95
2018 Feb	238.2103	[215.6829, 260.7378]80	[203.7575, 272.6632]95
2018 Mar	259.8648	[234.5976, 285.1321]80	[221.2219, 298.5078]95
2018 Apr	249.9134	[224.9679, 274.8590]80	[211.7625, 288.0643]95
2018 May	260.7056	[234.0279, 287.3833]80	[219.9056, 301.5056]95
2018 Jun	250.6406	[224.3801, 276.9012]80	[210.4786, 290.8027]95
2018 Jul	260.4818	[232.5695, 288.3941]80	[217.7937, 303.1699]95
2018 Aug	263.7202	[234.8480, 292.5925]80	[219.5639, 307.8766]95
2018 Sep	261.3672	[232.1593, 290.5750]80	[216.6976, 306.0367]95
2018 Oct	267.6420	[237.1392, 298.1449]80	[220.9920, 314.2921]95
2018 Nov	280.6423	[248.0488, 313.2358]80	[230.7949, 330.4897]95
2018 Dec	365.2664	[322.0688, 408.4641]80	[299.2014, 431.3315]95

```

MAM_fc %>%
  autoplot() +
  autolayer(myseries %>% filter_index("Dec 2015" ~ .)) +
  labs(
    title = "MAM Point Forecast and Point Interval 2 years after 2016 Dec",
    y = "Turnover [million $AUD]",
    x = "Month [M]"
  )

```

MAM Point Forecast and Point Interval 2 years after 2016 Dec



Best ARIMA Model

```
ARIMA1300013 <- trained_arimamodel %>%
  select(ARIMA1300013)
```

```
trained_arimamodel %>%
  accuracy() %>%
  arrange(MASE)
```


```
## # A tibble: 10 x 12
##   State   Industry .model .type      ME  RMSE  MAE      MPE  MAPE  MASE  RMSSE
##   <chr>   <chr>   <chr> <chr>   <dbl> <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl>
## 1 South A~ Other r~ ARIMA~ Trai~ -0.0109 4.95 3.56 0.0193 2.79 0.423 0.456
## 2 South A~ Other r~ ARIMA~ Trai~ 0.0341 5.04 3.58 0.103 2.85 0.425 0.464
## 3 South A~ Other r~ ARIMA~ Trai~ 0.0343 5.04 3.58 0.104 2.85 0.425 0.464
## 4 South A~ Other r~ ARIMA~ Trai~ 0.0109 5.05 3.58 0.0543 2.77 0.425 0.465
## 5 South A~ Other r~ ARIMA~ Trai~ 0.0104 5.05 3.58 0.0529 2.77 0.425 0.465
## 6 South A~ Other r~ ARIMA~ Trai~ -0.156 4.99 3.60 -0.137 2.82 0.427 0.460
## 7 South A~ Other r~ ARIMA~ Trai~ -0.0929 5.08 3.64 -0.0816 2.81 0.432 0.468
## 8 South A~ Other r~ ARIMA~ Trai~ -0.0888 5.09 3.65 -0.0785 2.82 0.433 0.468
## 9 South A~ Other r~ ARIMA~ Trai~ 0.0125 5.06 3.65 0.00296 2.92 0.433 0.466
## 10 South A~ Other r~ ARIMA~ Trai~ 0.0652 5.09 3.68 0.0462 2.94 0.436 0.469
## # ... with 1 more variable: ACF1 <dbl>
```

```
ARIMA1300013 %>%
  report()
```

```
## Series: Turnover
## Model: ARIMA(3,0,0)(0,1,3)[12] w/ drift
## Transformation: box_cox(Turnover, lambda)
##
## Coefficients:
##          ar1      ar2      ar3      sma1      sma2      sma3  constant
##          0.5304  0.1333  0.2926 -0.8467 -0.1221 -0.0312   0.0037
## s.e.      0.0479  0.0543  0.0491  0.1039  0.0665  0.0566   0.0002
##
## sigma^2 estimated as 0.002873:  log likelihood=594.55
## AIC=-1173.09  AICc=-1172.73  BIC=-1141.06
```

The best ARIMA model was model ARIMA1300013 since it had the lowest AIC value and during the testing set it also had a low MASE score which was tied for second place with four other models and it was only 0.002 off first place.

```
ARIMA1300013 %>%
  gg_tsresiduals() + labs(title = "ARIMA Innotvation Residuals Plot")
```



```
ARIMA1300013 %>%
  augment() %>%
  features(.innov, ljung_box, lag = 12, dof = 6) %>%
  kable(caption = "ARIMA1300013 Ljung Box Test")
```

Table 12: ARIMA1300013 Ljung Box Test

.model	lb_stat	lb_pvalue
ARIMA1300013	37.98809	1.1e-06

Similar to the ETS model, the innovation residuals plot and its corresponding ACF plot for the ARIMA model also shows that there might still be some correlation that was missed. Which means that model is not perfect and will not provide the most accurate forecasts. With that being said, the the true value will still land with in its forecast interval.

```
ARIMA1300013_fc <- ARIMA1300013 %>%
  forecast(h = "2 years")

ARIMA1300013_fc %>%
  hilo(level = c(80, 95)) %>%
  select(-.model, -Turnover) %>%
  kable(caption = "ARIMA Point Forecast and Point Interval (80% & 95%)")
```


Table 13: ARIMA Point Forecast and Point Interval (80% & 95%)

Month	.mean	80%	95%
2017 Jan	248.8701	[238.0793, 259.8311]	[232.6080, 265.8853]
2017 Feb	241.7388	[229.8604, 253.8299]	[223.8728, 260.5463]
2017 Mar	261.2925	[247.7188, 275.1236]	[240.8960, 282.8277]
2017 Apr	254.5790	[239.8774, 269.5913]	[232.5313, 278.0010]
2017 May	264.8888	[248.5508, 281.5961]	[240.4197, 290.9915]
2017 Jun	257.8743	[241.0990, 275.0499]	[232.7792, 284.7407]
2017 Jul	268.9679	[250.6500, 287.7442]	[241.5934, 298.3700]
2017 Aug	272.6887	[253.3403, 292.5421]	[243.8023, 303.8090]
2017 Sep	272.2290	[252.2063, 292.7939]	[242.3623, 304.4944]
2017 Oct	278.4553	[257.3419, 300.1587]	[246.9861, 312.5345]
2017 Nov	292.6486	[269.9036, 316.0455]	[258.7695, 329.4119]
2017 Dec	371.0949	[342.0798, 400.9468]	[327.8831, 418.0087]
2018 Jan	269.1068	[246.4556, 292.4629]	[235.4414, 305.8910]
2018 Feb	259.3883	[236.8081, 282.6961]	[225.8614, 296.1344]
2018 Mar	282.4194	[257.3813, 308.2798]	[245.2633, 323.2134]
2018 Apr	274.2942	[249.2740, 300.1607]	[237.1971, 315.1356]
2018 May	285.8053	[259.2248, 313.3031]	[246.4188, 329.2504]
2018 Jun	277.3682	[250.9813, 304.6874]	[238.2970, 320.5645]
2018 Jul	287.8944	[260.0712, 316.7169]	[246.7178, 333.4925]
2018 Aug	291.9333	[263.2639, 321.6497]	[249.5269, 338.9722]
2018 Sep	290.9595	[261.9266, 321.0704]	[248.0385, 338.6502]
2018 Oct	298.9568	[268.7591, 330.2900]	[254.3327, 348.6057]
2018 Nov	313.8220	[281.8258, 347.0332]	[266.5553, 366.4647]
2018 Dec	400.2891	[359.7947, 442.3088]	[340.4521, 466.8748]

```
ARIMA1300013_fc %>%
  autoplot() +
  autolayer(myseries %>% filter_index("Dec 2015" ~ .)) +
  labs(
    title = "ARIMA1300013 Point Forecast and Point Interval 2 years after 2016 Dec",
    y = "Turnover [million $AUD]",
```

```
x = "Month [M]"
)
```

ARIMA1300013 Point Forecast and Point Interval 2 years after 2016 Dec



Comparison between ETS and ARIMA Models

```
bind_rows(
  ARIMA1300013 %>% accuracy(),
  MAM %>% accuracy(),
  ARIMA1300013 %>% forecast(h = 24) %>% accuracy(myseries),
  MAM %>% forecast(h = 24) %>% accuracy(myseries)
) %>%
  select(-ME, -MPE, -ACF1) %>%
  arrange(MASE) %>%
  kable(caption = "Comparison between ETS and ARIMA Models")
```

Table 14: Comparison between ETS and ARIMA Models

.model	.type	RMSE	MAE	MAPE	MASE	RMSSE
ARIMA1300013	Training	5.048520	3.582733	2.765884	0.4250913	0.4647777
MAM	Training	5.119750	3.631144	2.852607	0.4308353	0.4713354
MAM	Test	8.277969	6.249165	2.226650	0.7414636	0.7620879
ARIMA1300013	Test	22.341353	18.350703	6.745479	2.1773114	2.0567937

After forecasting 2 years after the end of the training data, the accuracy of the ETS model triumphs the

ARIMA model. The ETS model has a MASE value of 0.7415 while the ARIMA model had 2.1773.

Forecasting


```
fit %>%
  select(MAM) %>%
  forecast(h = "2 years") %>%
  hilo(level = 80) %>%
  select(-Turnover) %>%
  kable(
    caption = "MAM Point Forecast and Point Interval 2 years after 2018 Dec"
  )
```

Table 15: MAM Point Forecast and Point Interval 2 years after 2018 Dec

.model	Month	.mean	80%
MAM	2019 Jan	252.3784	[240.7473, 264.0095]
MAM	2019 Feb	239.3960	[227.3504, 251.4416]
MAM	2019 Mar	261.3224	[247.1564, 275.4884]
MAM	2019 Apr	252.1888	[237.6043, 266.7733]
MAM	2019 May	262.6367	[246.5548, 278.7186]
MAM	2019 Jun	252.6429	[236.3610, 268.9248]
MAM	2019 Jul	261.9044	[244.2257, 279.5831]
MAM	2019 Aug	265.1500	[246.4793, 283.8206]
MAM	2019 Sep	262.2458	[243.0473, 281.4443]
MAM	2019 Oct	269.3486	[248.9072, 289.7899]
MAM	2019 Nov	283.5678	[261.3140, 305.8216]
MAM	2019 Dec	367.2047	[337.4701, 396.9393]
MAM	2020 Jan	257.3701	[235.9069, 278.8333]
MAM	2020 Feb	244.1232	[223.1925, 265.0538]
MAM	2020 Mar	266.4740	[243.0204, 289.9276]
MAM	2020 Apr	257.1522	[233.9497, 280.3547]
MAM	2020 May	267.7973	[243.0569, 292.5377]
MAM	2020 Jun	257.5990	[233.2593, 281.9388]
MAM	2020 Jul	267.0338	[241.2547, 292.8129]
MAM	2020 Aug	270.3345	[243.6948, 296.9742]
MAM	2020 Sep	267.3652	[240.4940, 294.2364]
MAM	2020 Oct	274.5981	[246.4731, 302.7231]
MAM	2020 Nov	289.0855	[258.9334, 319.2375]
MAM	2020 Dec	374.3382	[334.6046, 414.0717]

```
fit %>%
  select(MAM) %>%
  forecast(h = "2 years") %>%
  autoplot() +
  autolayer(myseries %>% filter_index("Dec 2015" ~ .)) +
  labs(
    title = "MAM Point Forecast and Point Interval 2 years after 2018 Dec",
    y = "Turnover [million $AUD]",
    x = "Month [M]"
  )
```

MAM Point Forecast and Point Interval 2 years after 2018 Dec



```
fit %>%
  select(ARIMA1300013) %>%
  forecast(h = "2 years") %>%
  hilo(level = 80) %>%
  select(-Turnover) %>%
  kable(caption = "ARIMA1300013 Point Forecast and Point Interval 2 years after 2018 Dec")
```

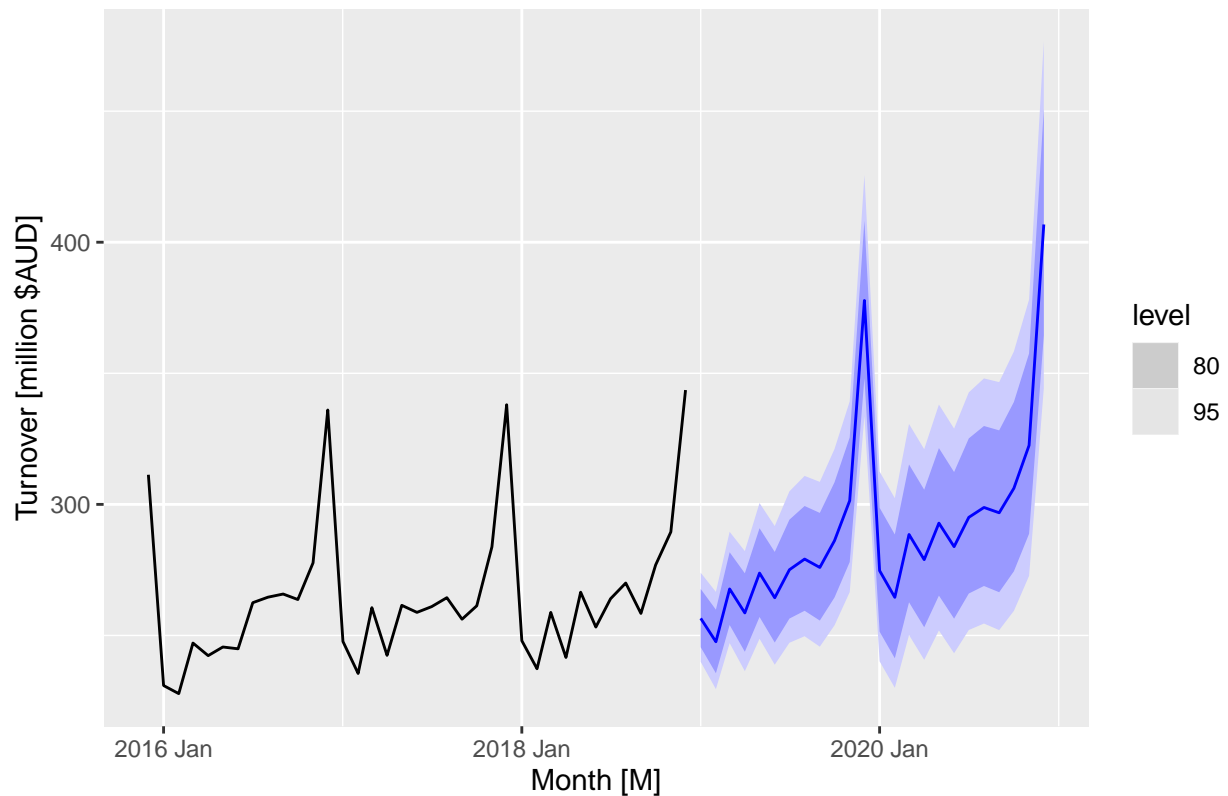
Table 16: ARIMA1300013 Point Forecast and Point Interval 2 years after 2018 Dec

.model	Month	.mean	80%
ARIMA1300013	2019 Jan	256.5240	[245.5063, 267.7138]80
ARIMA1300013	2019 Feb	247.6041	[235.5939, 259.8266]80
ARIMA1300013	2019 Mar	267.7421	[254.0010, 281.7406]80
ARIMA1300013	2019 Apr	258.5913	[243.7867, 273.7060]80
ARIMA1300013	2019 May	273.7926	[257.0649, 290.8947]80
ARIMA1300013	2019 Jun	264.3879	[247.3187, 281.8614]80
ARIMA1300013	2019 Jul	275.0657	[256.4380, 294.1568]80
ARIMA1300013	2019 Aug	279.1386	[259.4179, 299.3716]80
ARIMA1300013	2019 Sep	275.9387	[255.6847, 296.7401]80
ARIMA1300013	2019 Oct	286.1578	[264.4960, 308.4239]80
ARIMA1300013	2019 Nov	301.4412	[278.0206, 325.5329]80
ARIMA1300013	2019 Dec	377.7701	[348.1696, 408.2261]80
ARIMA1300013	2020 Jan	274.7090	[251.4418, 298.7049]80
ARIMA1300013	2020 Feb	264.5326	[241.3076, 288.5126]80
ARIMA1300013	2020 Mar	288.5125	[262.6653, 315.2177]80

.model	Month	.mean	80%
ARIMA1300013	2020 Apr	278.9207	[253.1364, 305.5893]80
ARIMA1300013	2020 May	292.8468	[265.2062, 321.4560]80
ARIMA1300013	2020 Jun	283.8716	[256.4086, 312.3219]80
ARIMA1300013	2020 Jul	295.0499	[265.9995, 325.1641]80
ARIMA1300013	2020 Aug	298.8241	[268.8690, 329.8968]80
ARIMA1300013	2020 Sep	296.7894	[266.4983, 328.2320]80
ARIMA1300013	2020 Oct	306.1597	[274.4841, 339.0566]80
ARIMA1300013	2020 Nov	322.5223	[288.7949, 357.5646]80
ARIMA1300013	2020 Dec	406.7703	[364.4557, 450.7258]80


```
fit %>%
  select(ARIMA1300013) %>%
  forecast(h = "2 years") %>%
  autoplot() +
  autolayer(myseries %>% filter_index("Dec 2015" ~ .)) +
  labs(
    title = "ARIMA1300013 Point Forecast and Point Interval 2 years after 2018 Dec",
    y = "Turnover [million $AUD]",
    x = "Month [M]"
  )
)
```

ARIMA1300013 Point Forecast and Point Interval 2 years after 2018 Dec



Point Forecast and Point Interval 2 years after 2018 Dec

MAM and ARIMA1300013




```
fit %>%
  forecast(h = "2 years") %>%
  accuracy(true_data)
```

```
## # A tibble: 2 x 10
##   .model      .type    ME  RMSE  MAE  MPE  MAPE  MASE  RMSSE  ACF1
##   <chr>      <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 ARIMA1300013 Test    2.11  10.7  8.76  0.747  2.94  1.06  0.998 -0.0311
## 2 MAM        Test   20.9  24.6  21.7  7.06  7.29  2.62  2.30  0.222
```

```
fit %>%
  forecast(h = "2 years") %>%
  autoplot() +
  autolayer(true_data %>% filter_index("Dec 2015" ~ .)) +
  labs(
    title = "Point Forecast and Point Interval 2 years after 2018 Dec",
    subtitle = "ABSDATA",
    y = "Turnover [million $AUD]",
    x = "Month [M]"
  )
```

Point Forecast and Point Interval 2 years after 2018 Dec ABSDATA



Checking the MASE of the forecasts and plots of the forecasts, we can tell that ARIMA1300013 did a decent job at predicting Turnover. MAM model consistently underestimated the mean Turnover but the true values were still within the prediction intervals.

Overall the both models were quite limiting since both the ETS and ARIMA models did not have its residuals as the Ljung Box test which is use to verify that variations between the fitted values and the real data are only due to randomness and not due to unaccounted for correlation. The models were also considerably simpler since a more complex model could potentially better fit the training data and testing data but it could also be negatively effected with its ability to predict future values. On the other hand a simpler model does not overfit the data which could allow it to better predict. So ultimately the best model would be somewhere in the middle.